Load balancing strategies in AIML networks

Costin Raiciu

Broadcom and Politehnica of Bucharest

Flow collisions

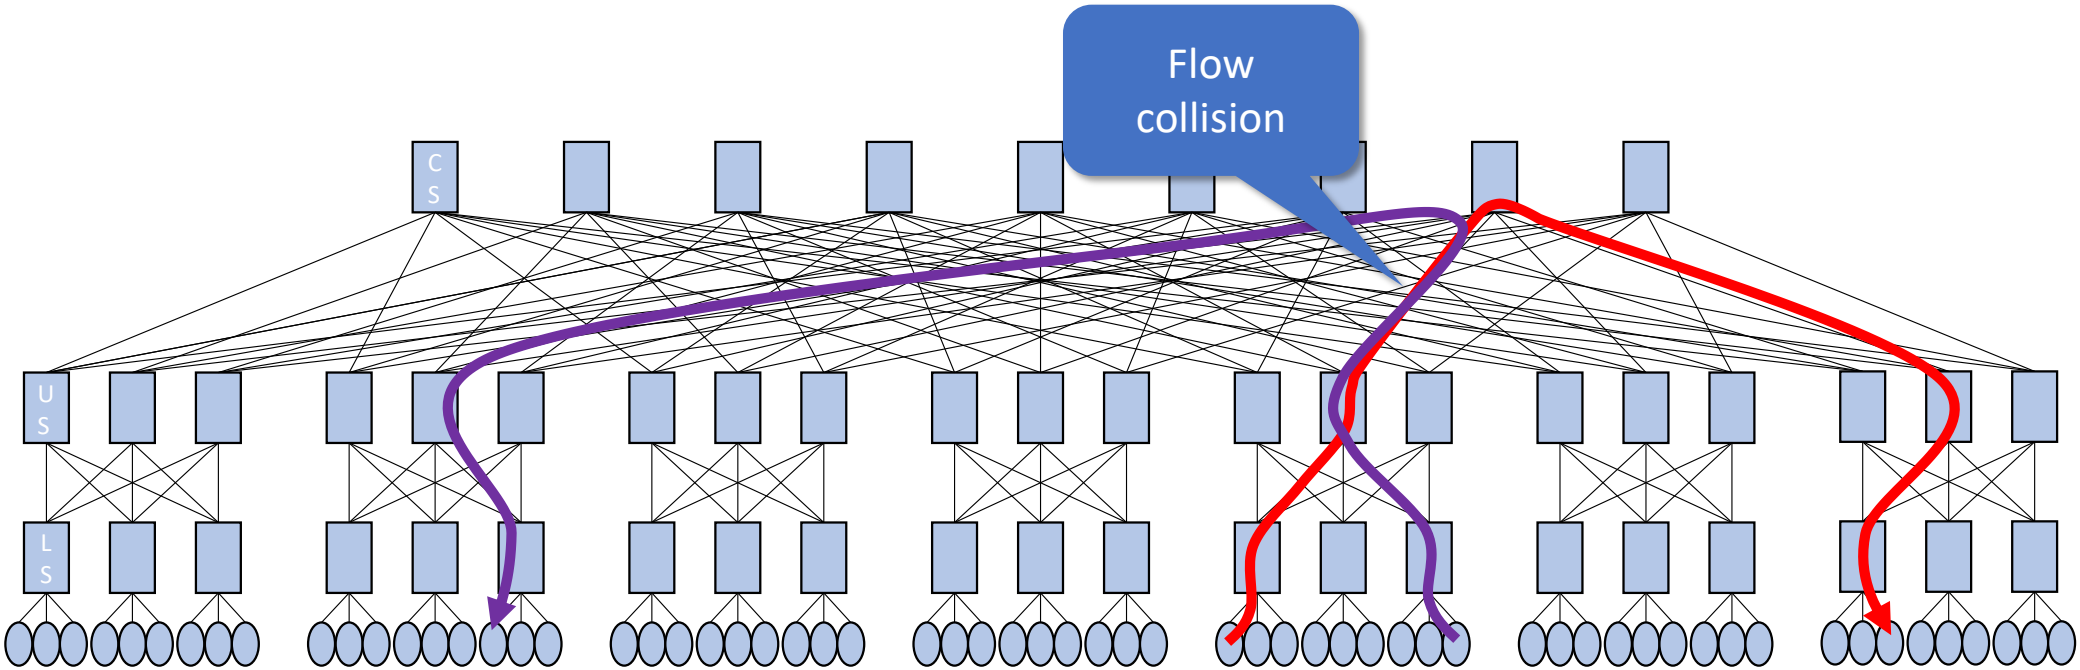

Why not use MPTCP for AIML networks?

- Need to use many paths for the common case of
  - Symmetric highly loaded networks + short flows.
  - Best way to load balance short flows is to use many paths.
  - Load balancing works well for long flows, not so well for shorter flows.
- But with many paths, minimum MPTCP total window is #paths.
  - E.g. 256 paths means min 256 packet window. This equals BDP at 800Gbps.
  - Congestion collapse in incast.
- Path state for MPTCP is quite costly.
  - CWND, flight_size, sequence numbers, etc. (tens of bytes).
  - 256 * 20 = 5KB per connection!

Packet spraying in AIML networks

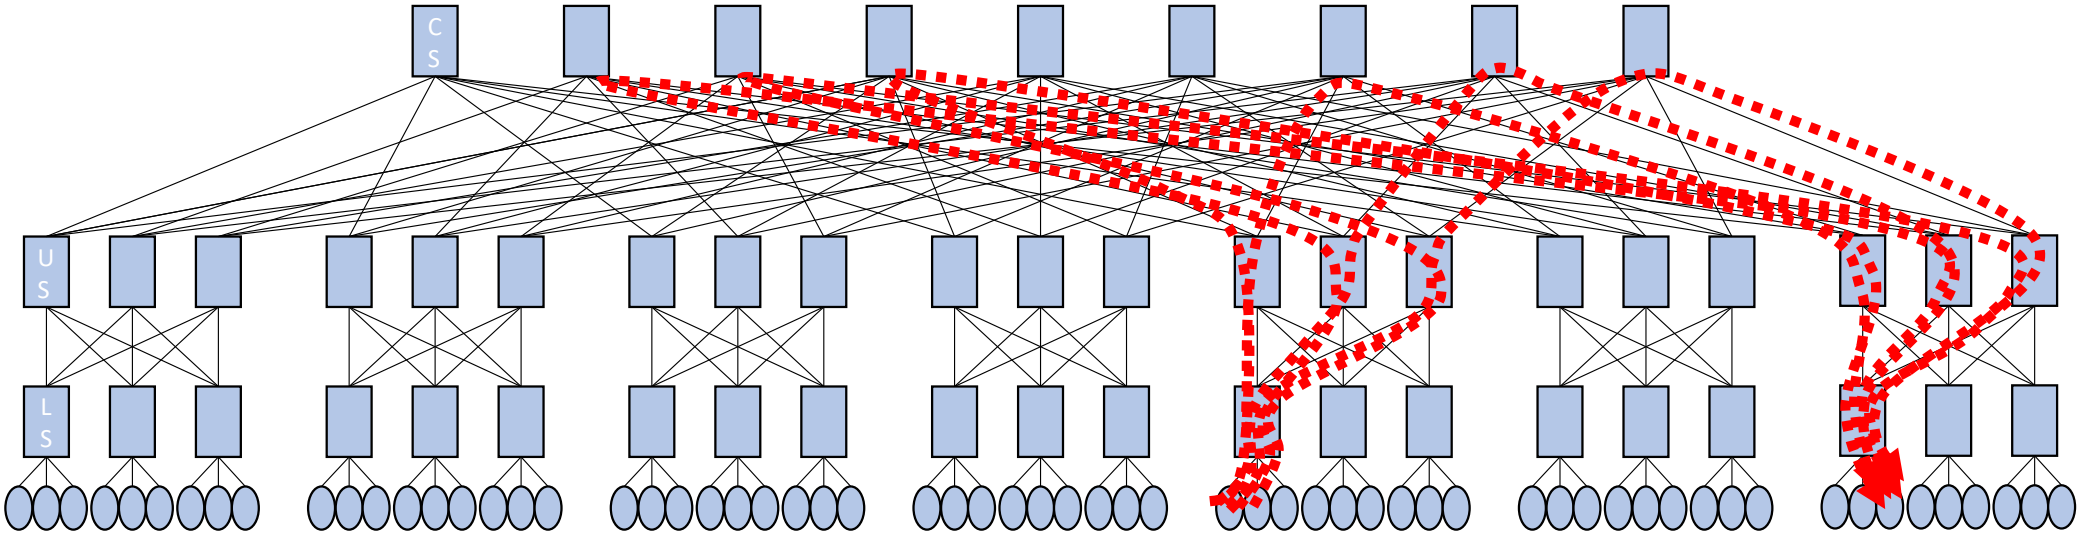

Congestion control with packet spraying

- Maintain a single congestion window that upper bounds flight size.
- Sender-driven congestion control (e.g. UET NSCC)
  - Targets sub-BDP standing queue at the bottleneck.
  - Use ECN and delay simultaneously.
    - Aggressive increase when queue ~ 0. Linear increase otherwise.
    - Multiplicative decrease when ECN mark & average delay above threshold.
  - Average delay across all paths.

- But how to load balance packets across paths?
  - Bad load balancing results in reducing CWND (across all paths).

End-to-end packet spraying via ECMP

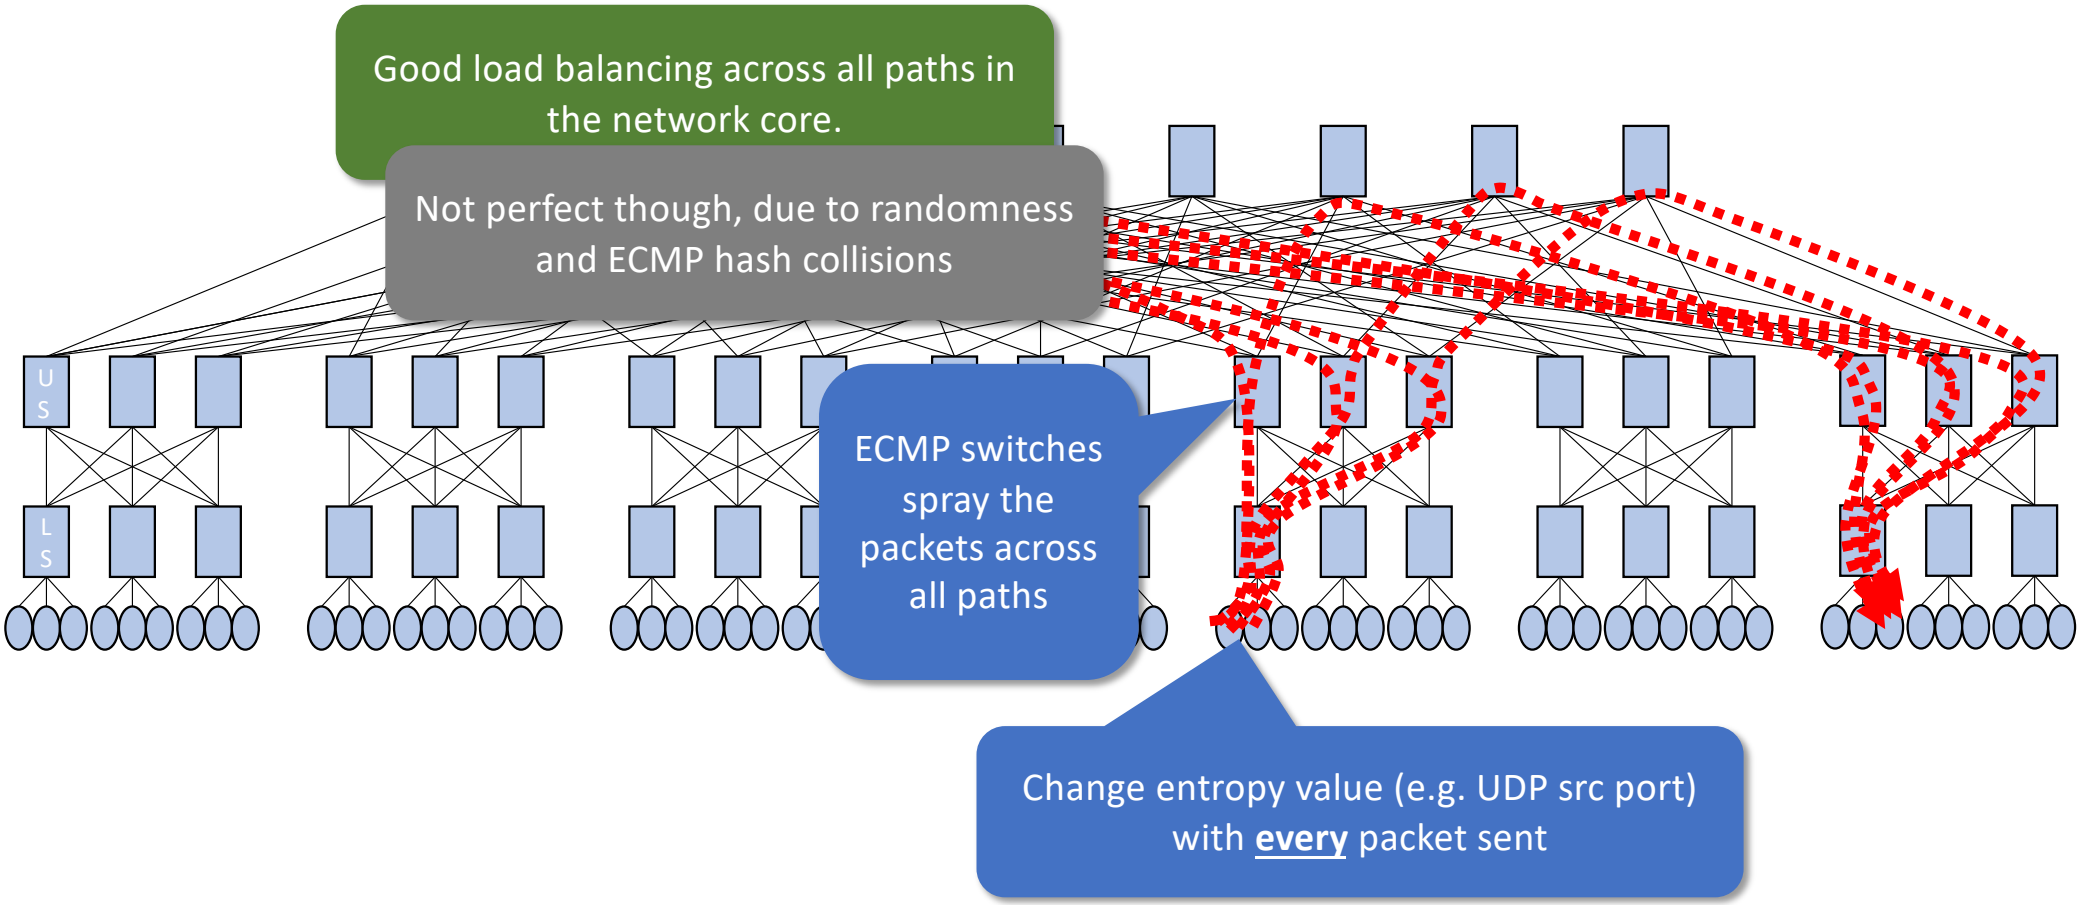

End-to-end packet spraying with source routing.

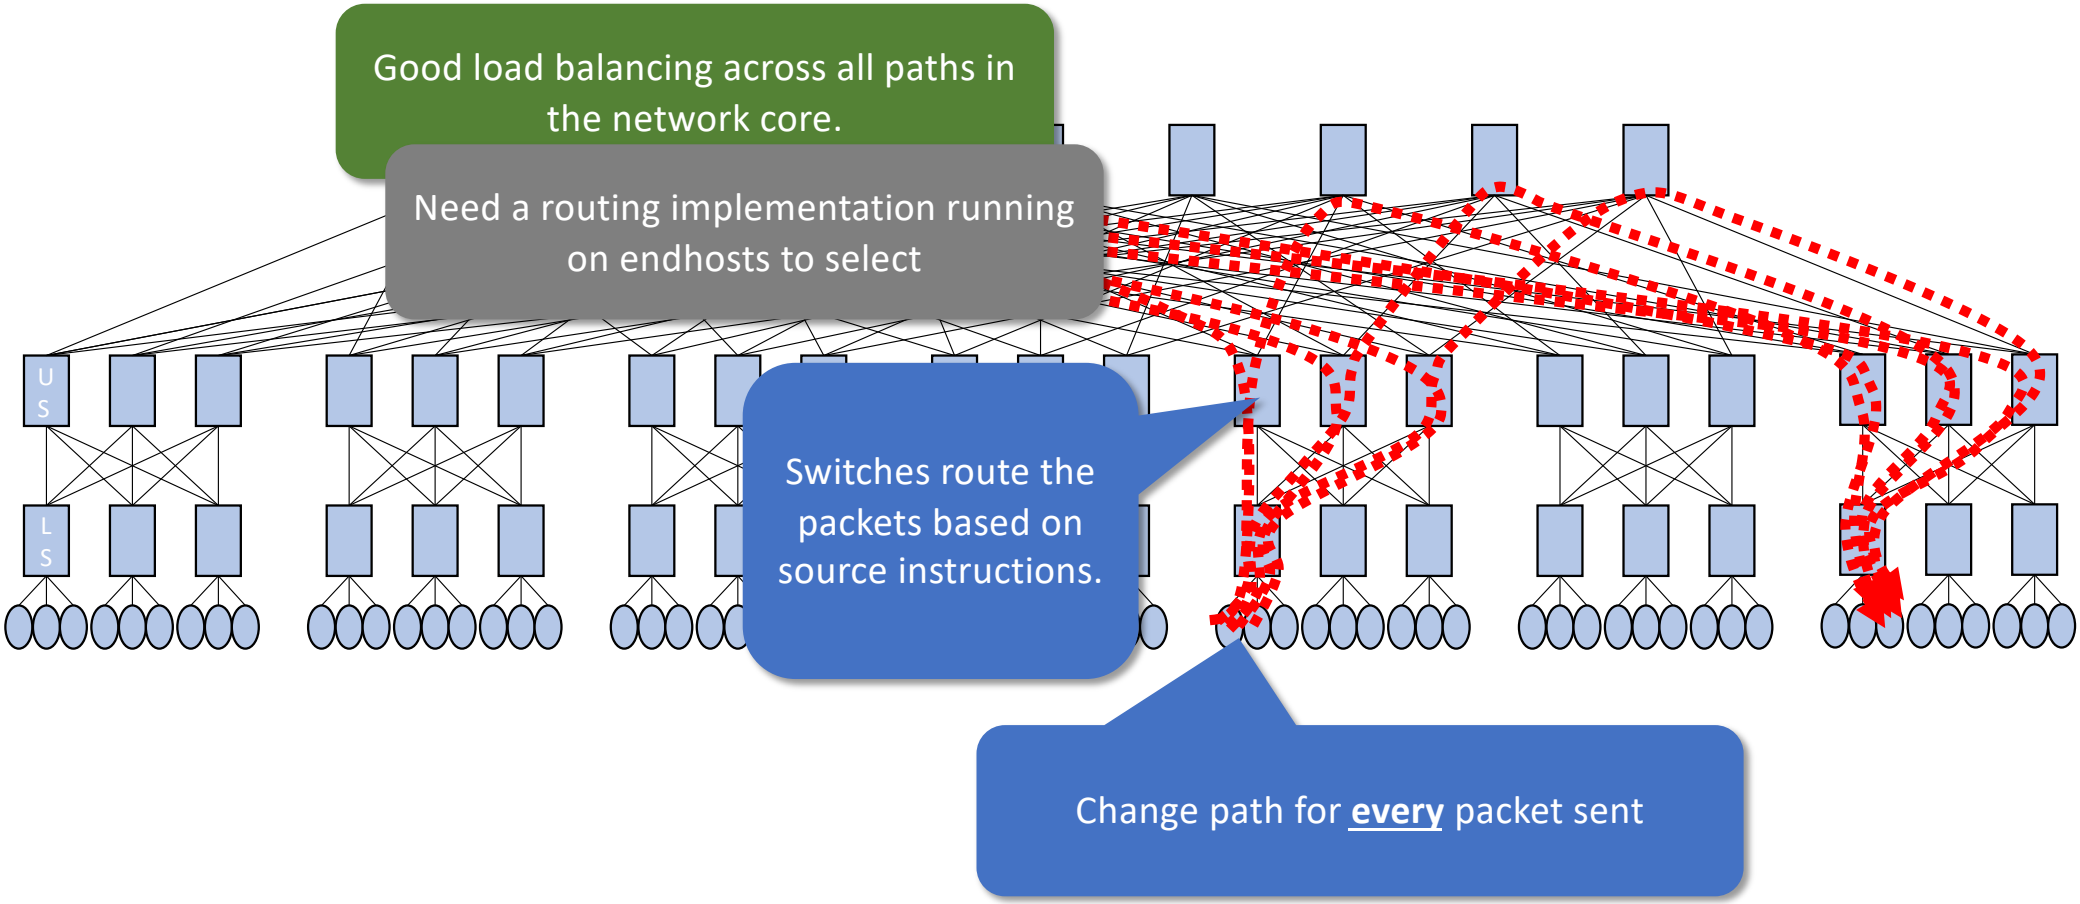

Switch packet spraying

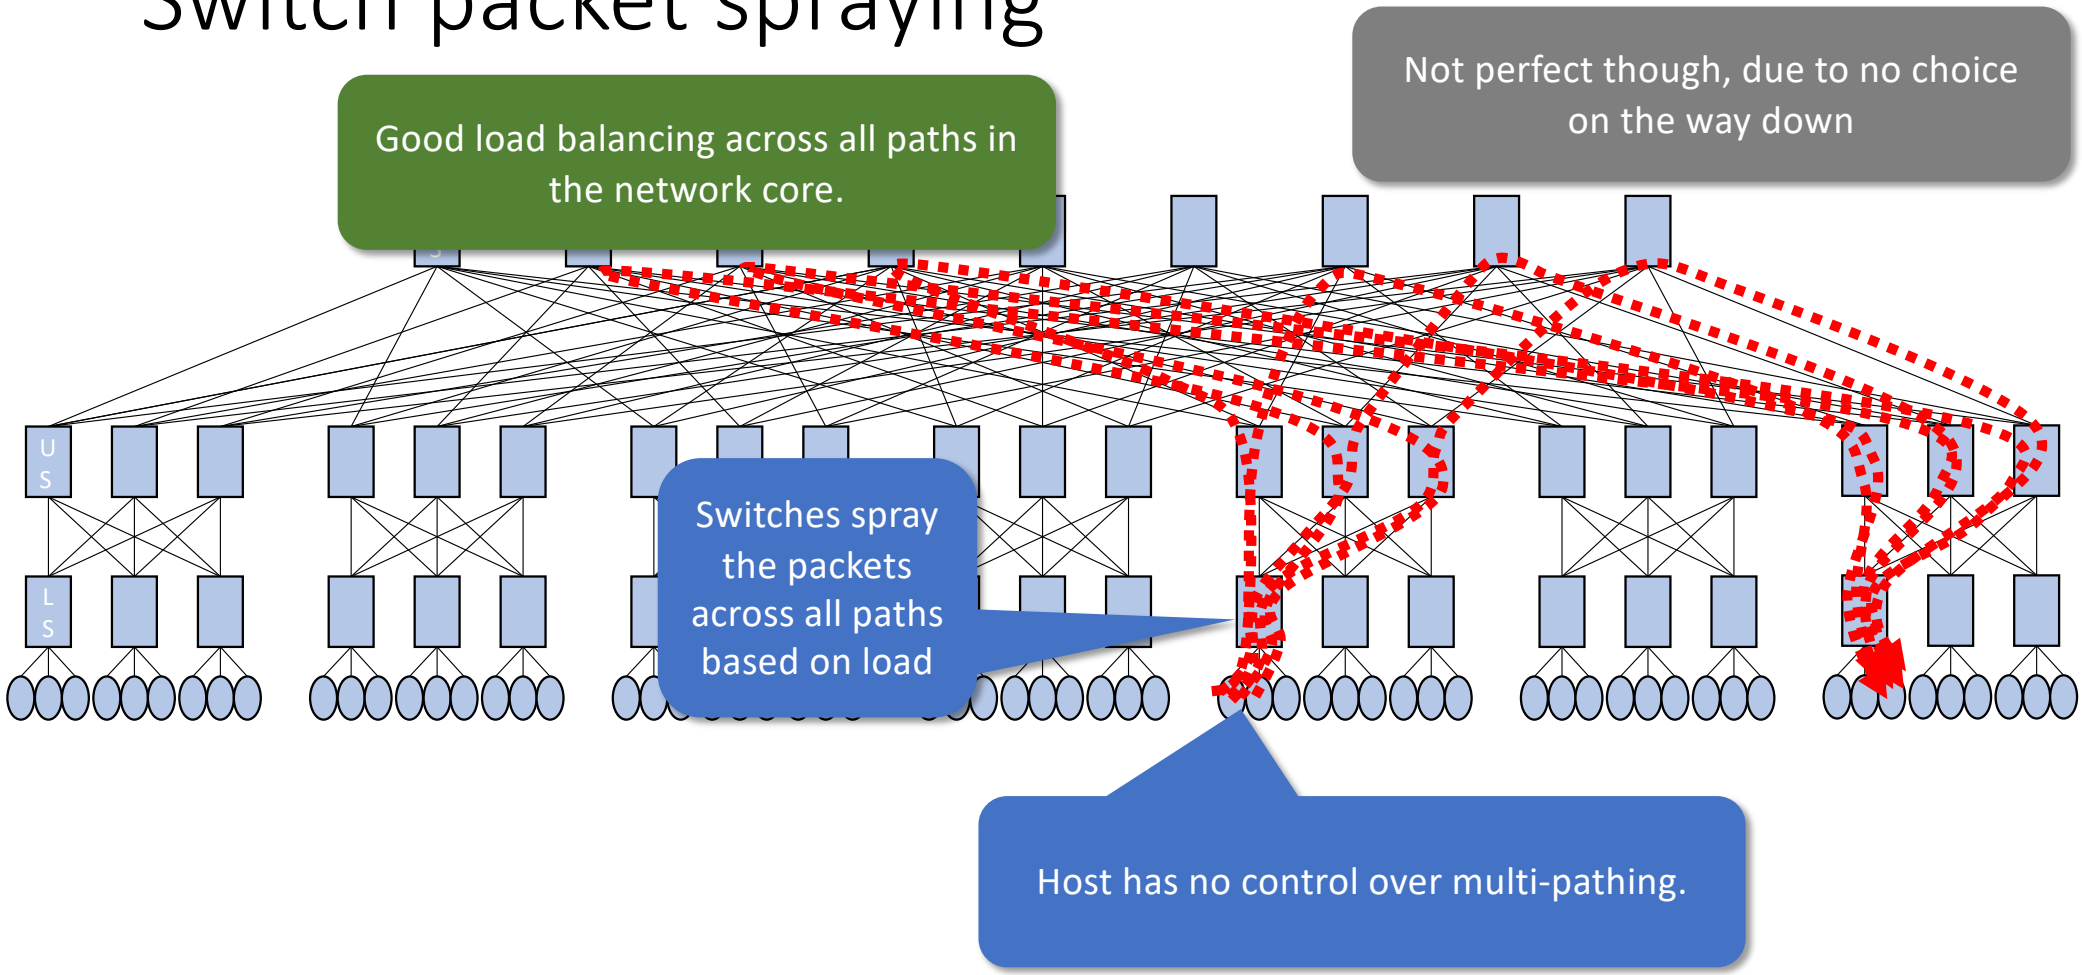

Two basic approaches for spraying

- Host-based spraying –
  - ECMP + standard routing protocol (e.g. BGP).
  - Source routing – requires an SDN routing protocol to compute paths and deliver them to hosts.

- Switch spraying – adaptive routing / dynamic load balancing.

What are the pros and cons of each?

Host based spraying
What is the best way to spray packets?

- **Simplest: oblivious load balancing**
  - Pick a random EV for each packet.
  - Works very well if network capacity is uniform.

- Bitmap load balancing (e.g. UET bitmap algorithm)
  - Per EV state - one or a few bits.
  - When ACK indicates ECN mark, increment EV state.
  - When EV is next to be picked but non-zero state, decrement state, skip.

- Recycled entropies (REPS):
  - Keep EV cache for which we got an ACK without ECN set.
  - Path selection: pick EV from cache if non-empty. Otherwise pick random EV.

Oblivious packe

When sprayed load balancing is imperfect, queues can still build.
Trimming prevents queue building.
Packet gets trimmed, NACKed,
RTX on a different less loaded path.

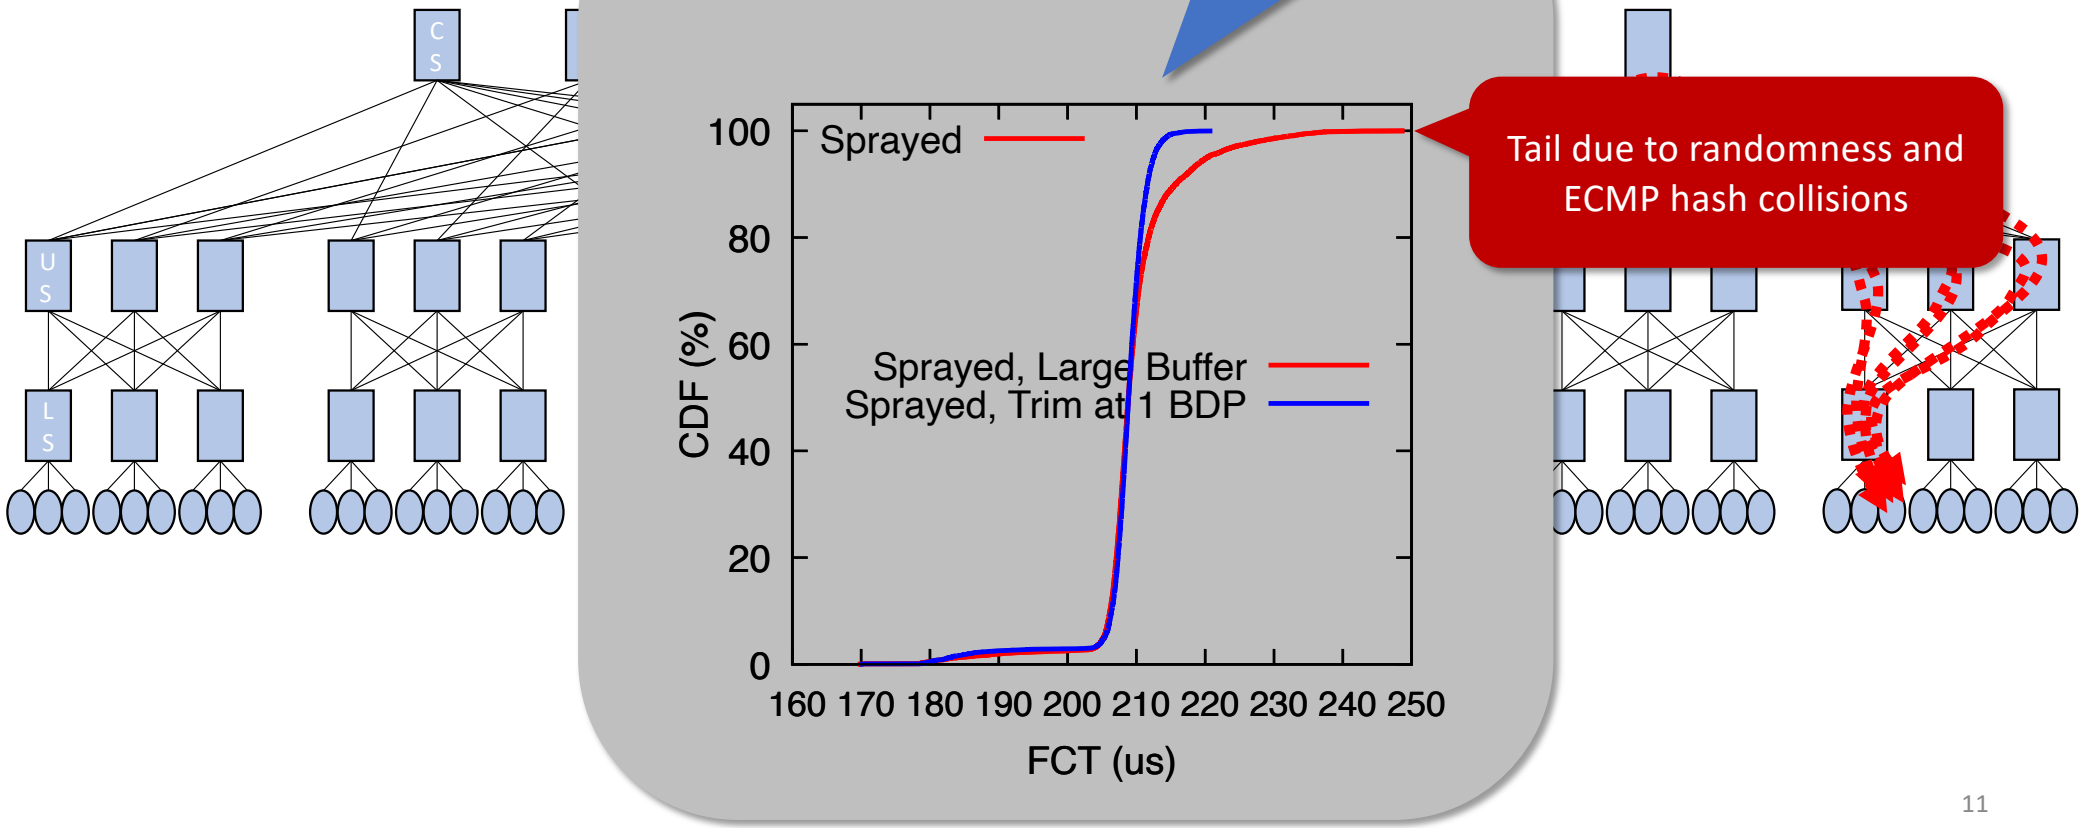

Tail due to randomness and ECMP hash collisions

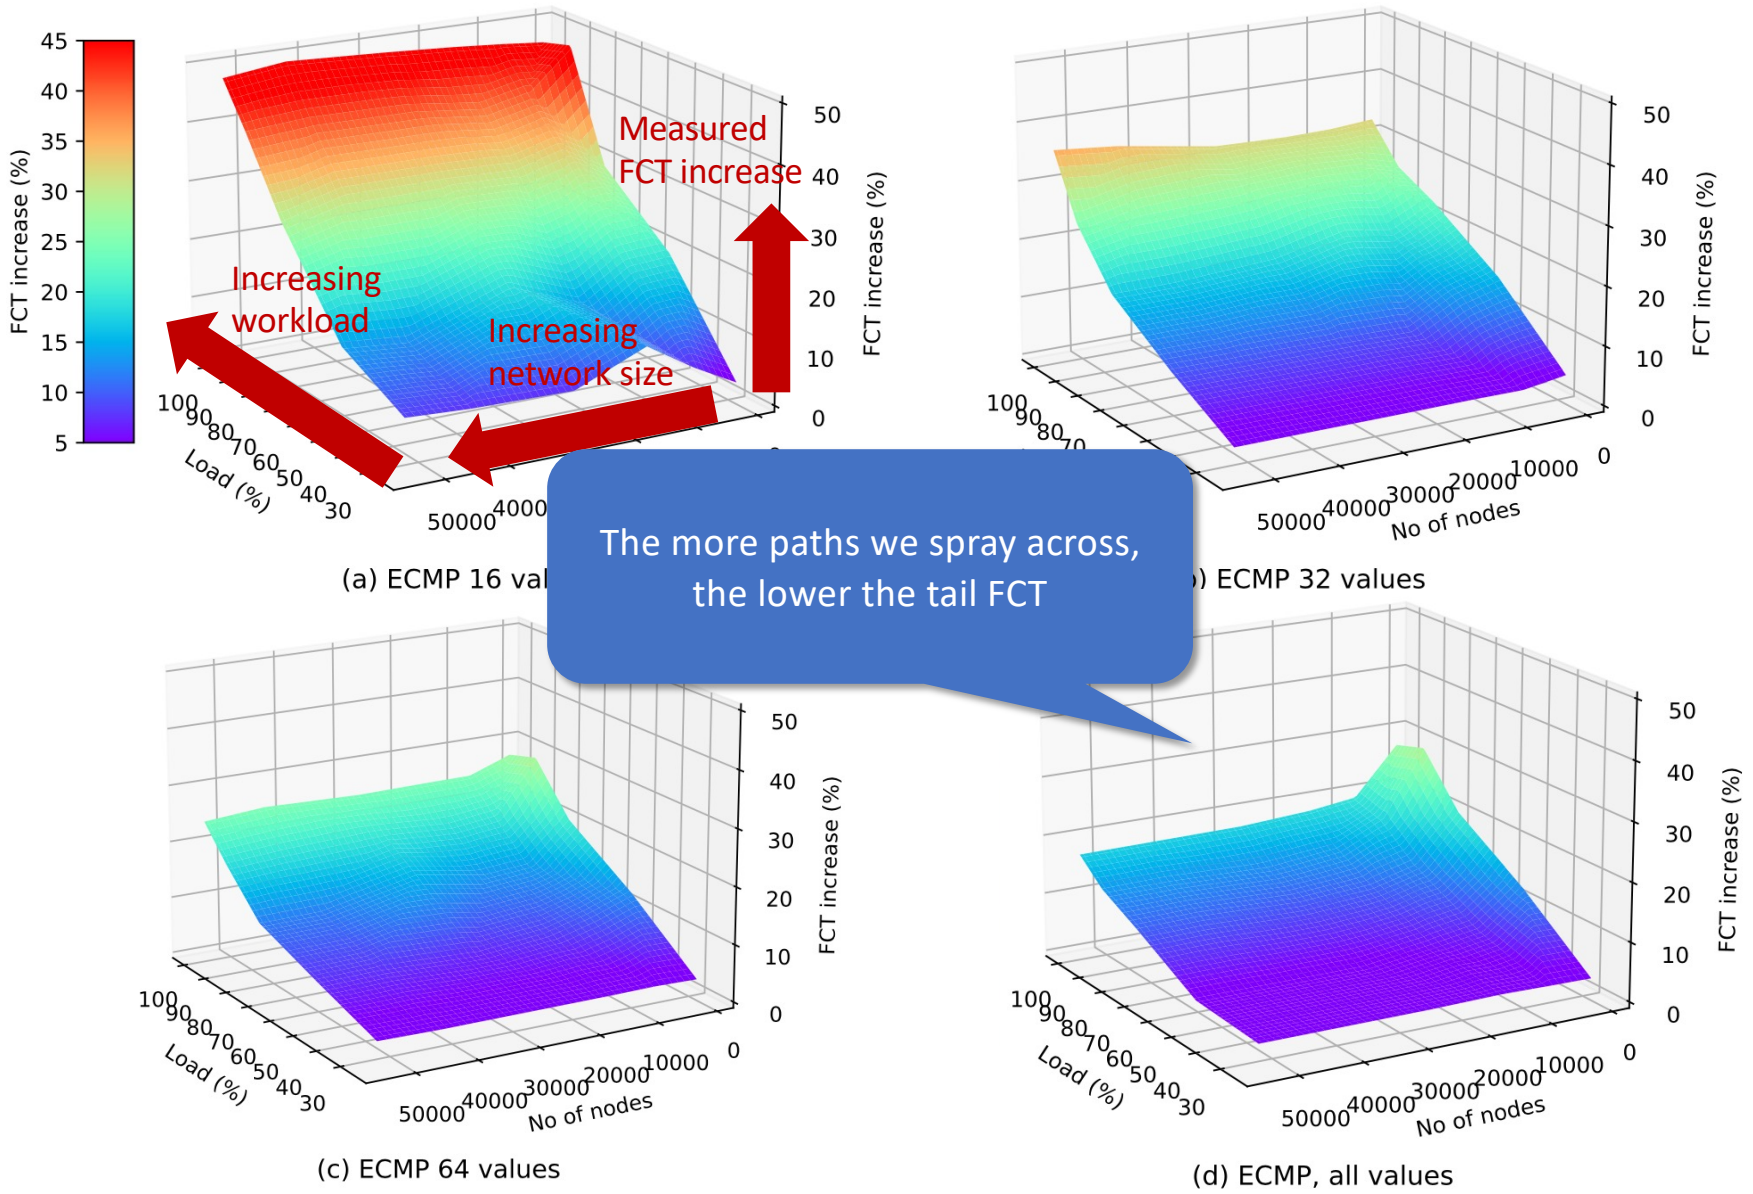
FCT increase (%)
Measured
FCT increase
Increasing
workload
Increasing
network size
Load (%)
No of nodes
The more paths we spray across,
the lower the tail FCT
(a) ECMP 16 val
) ECMP 32 values
(c) ECMP 64 values
(d) ECMP, all values

How does switch spray work?

- ECMP group => DLB group.
- During route lookup, routes in DLB group are consulted.
  - Contains all available paths towards destination.
  - Switch uses local information to decide which route (and associated egress port) to pick
  - Example metrics:
    - Queue length
    - Bandwidth utilization
    - PFC Port state.
    - Combination of the above possible.
- Works very well when path choice exists (e.g. going up the tree).
- Less well on the downward path / with asymmetries.
  - At the limit, behaves like oblivious endhost spraying.

Load balancing algorithms comparison

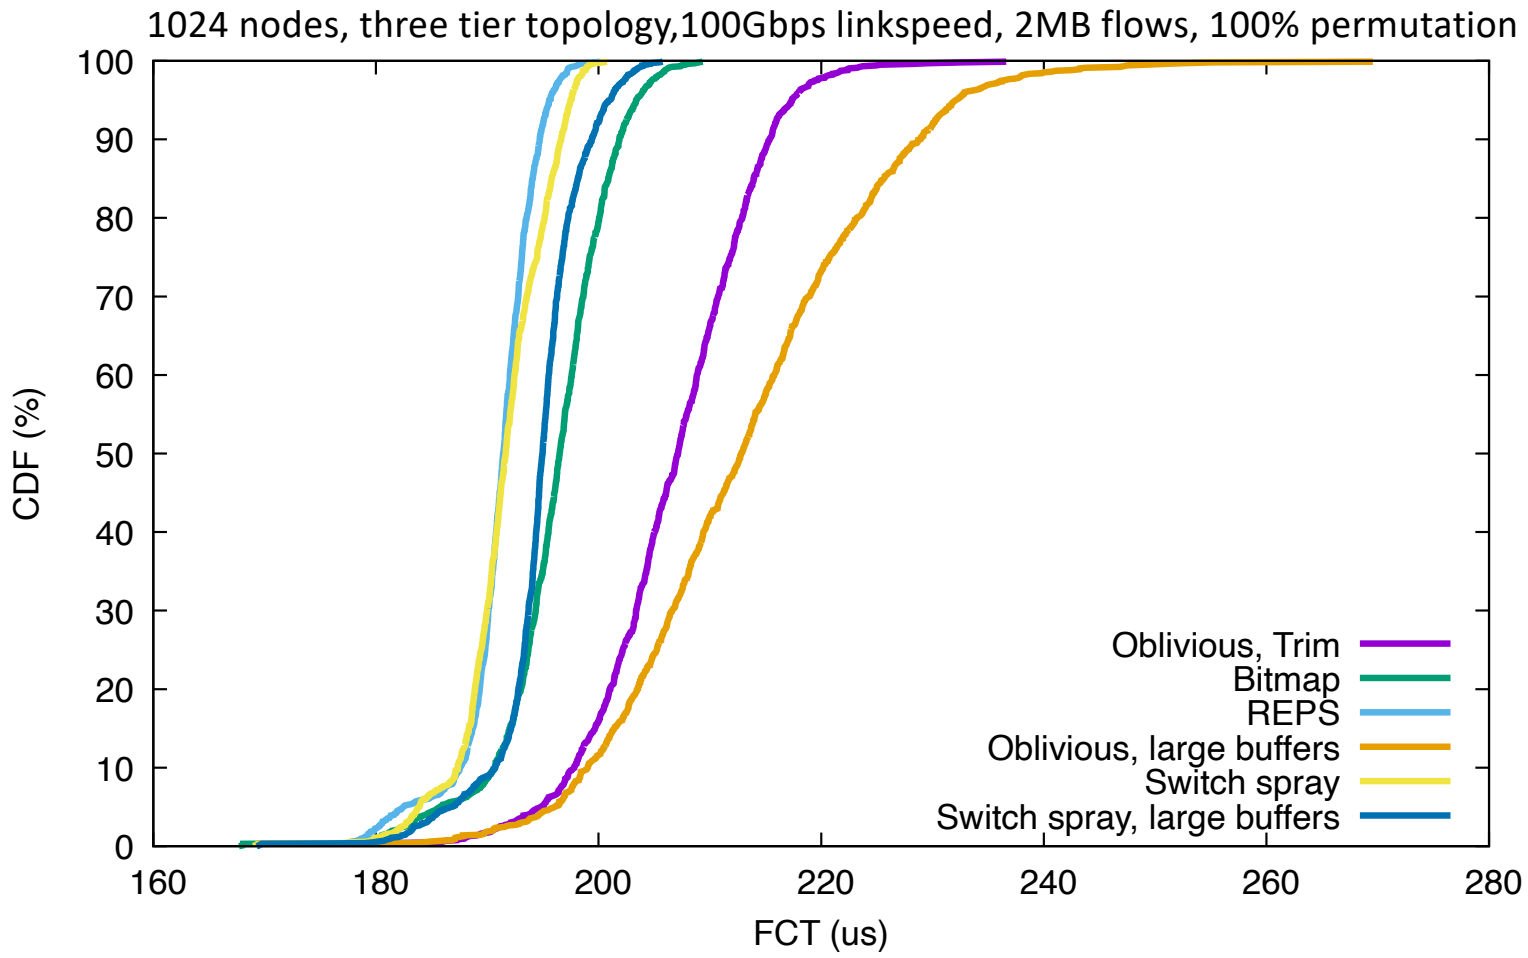

Load balancing with asymmetric capacity

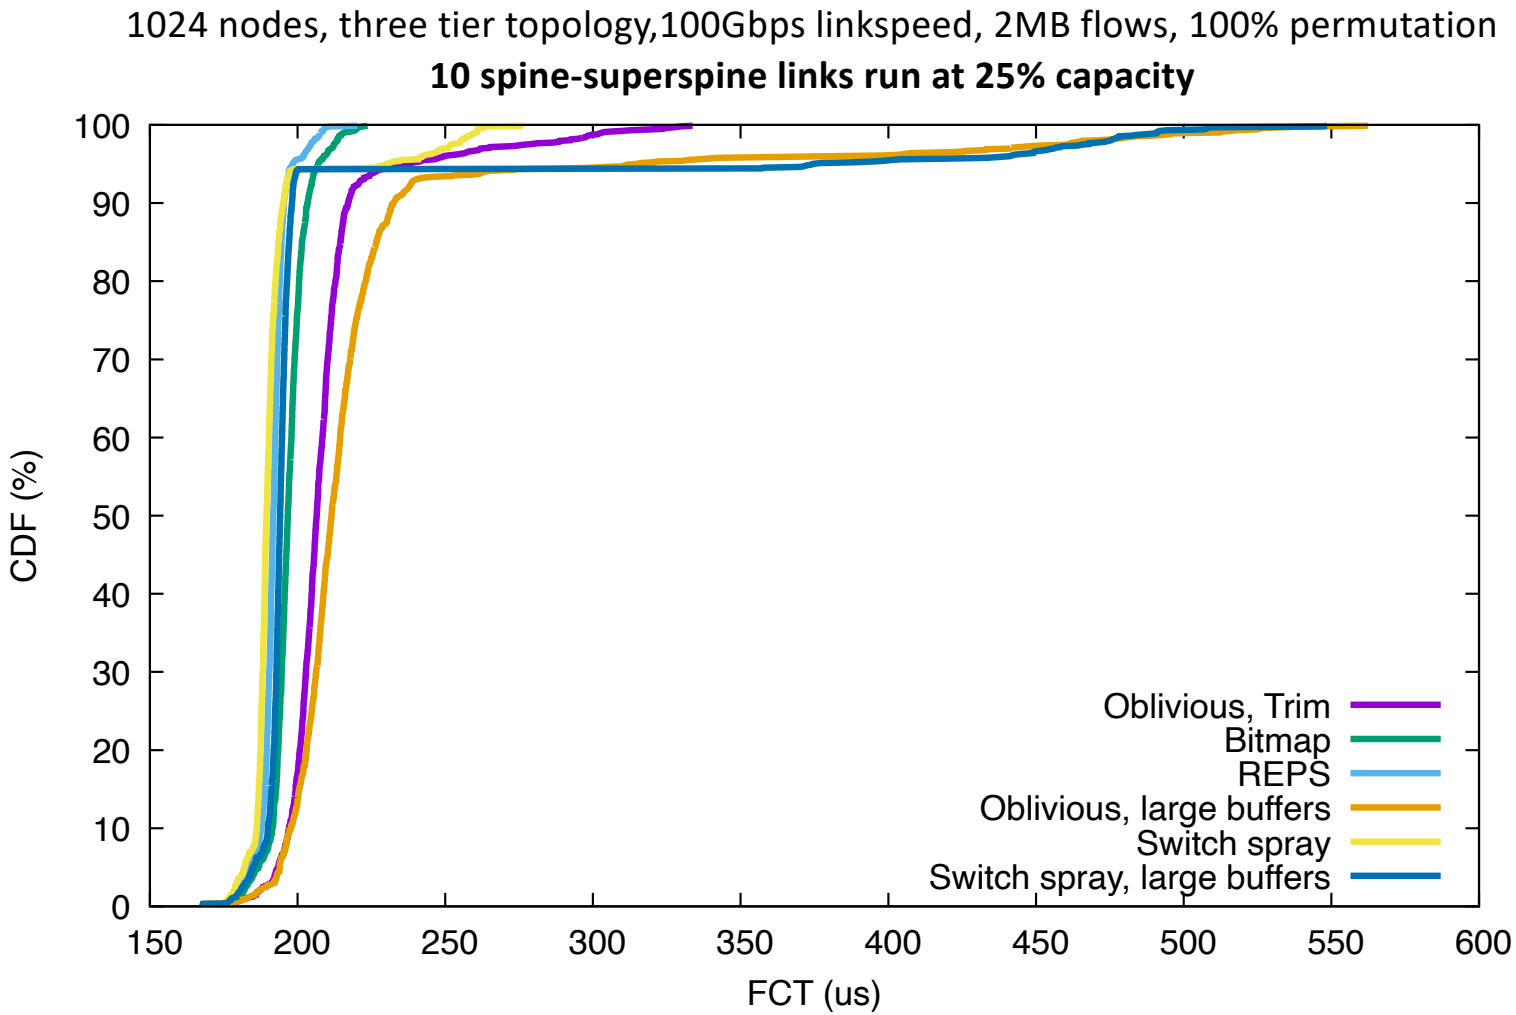

Summary

- Packet spraying enables exploring many/all paths.
- Endhost or switch spraying possible.
- All schemes work almost perfect when network is perfectly symmetric.
  - Switch spraying
    - Lower average buffer utilization than endhost load balancing.
    - Works very well when path choice available (e.g. link bundles).
  - Endhost spraying:
    - State-based schemes can achieve similar FCT to switch spraying.
    - But lead to higher queue utilization.
    - Even oblivious works quite well.
- When network is asymmetric, poor load balancing leads to large FCT.
  - Switch spray on its own struggle – need additional mechanisms.
  - State-based endhost load balancing copes fairly well.