

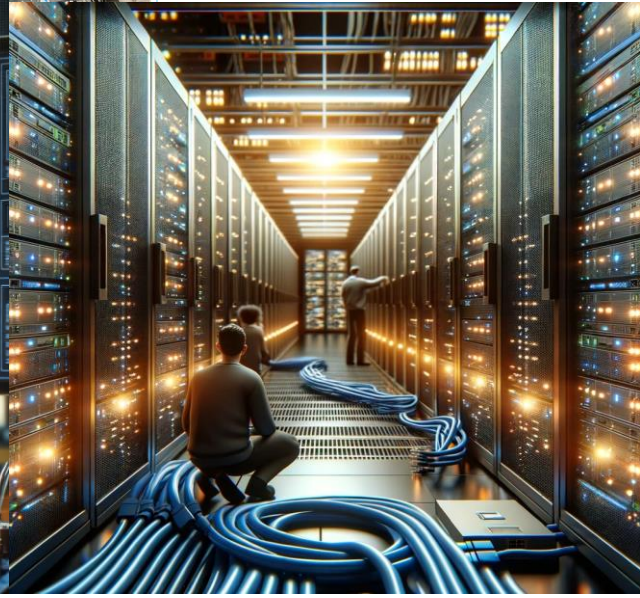

UCL

BROADCOM®

The current lossless abstraction

Everything sent shalt be received.

- If we lose something, e.g., a packet, we'll retransmit it.

The tail latency problem

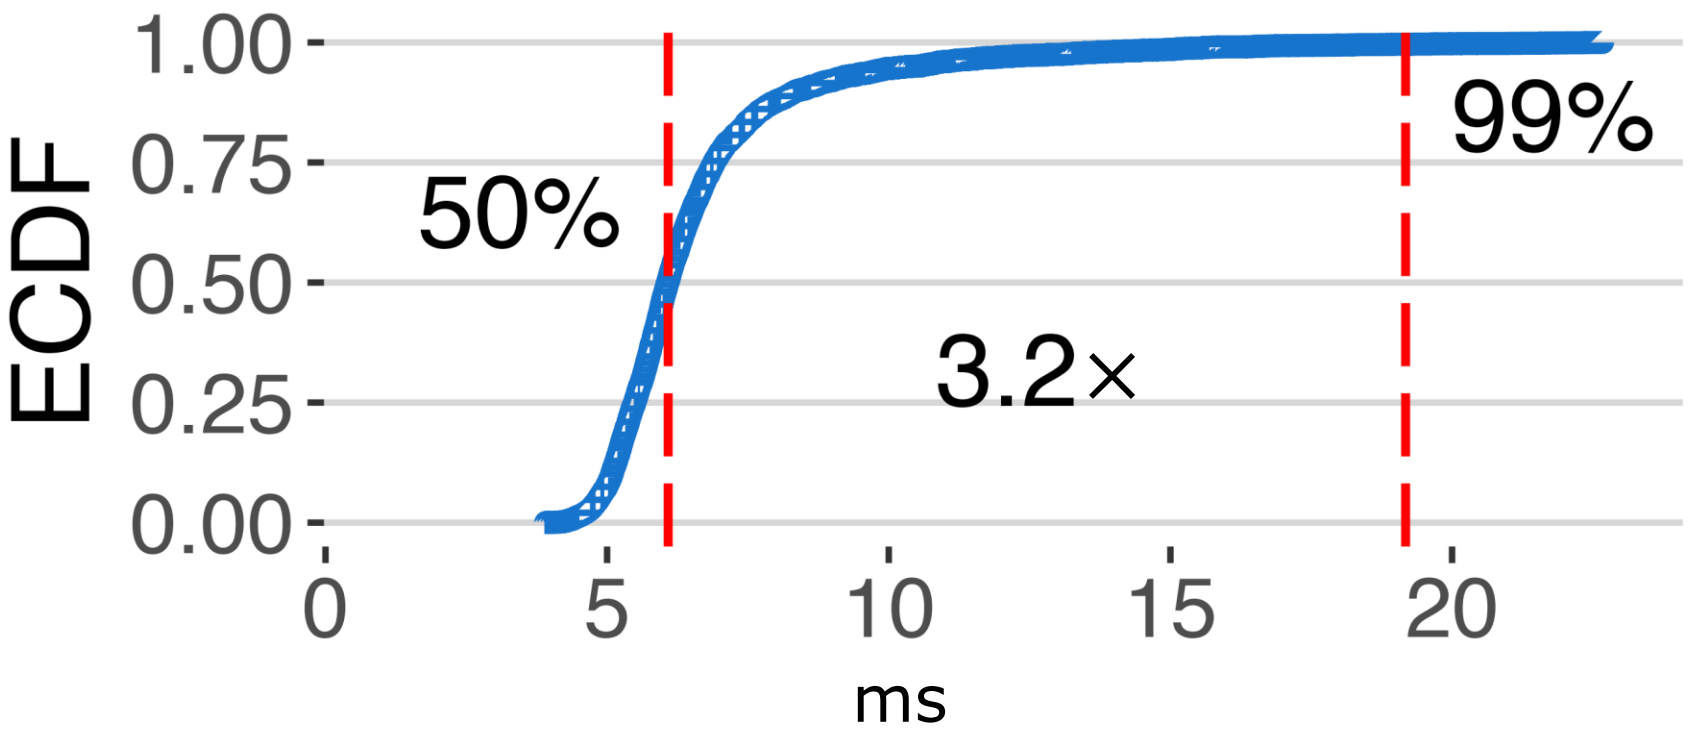

(OPTIREDUCE, NSDI 2025)

The tail latency problem

- Happens due to various reasons:
  - Packet drops due to congestion/corruption.
  - Queue buildups.
  - Accelerator failures.
  - Switch failures.
  - Link failures.
  - Straggler nodes.
  - …
- Exacerbated when training cross data centers.
- Mainly happens due to system events!

Embracing loss, a new paradigm

- We propose allowing the system flexibility in the gradient synchronization to mitigate the causes of the tail latency.
- To control the added loss, we need to formalize the allowed errors.

Example: Mitigating congestion with just-in-time compression (HotNets 2024)

- Congestion can increase latency of gradient synchronization and is not always avoidable.
- What if switches compress the data whenever needed?
- Can we do it with existing hardware switches?

Example: Mitigating compression with just-in-time compression (HotNets 2024)

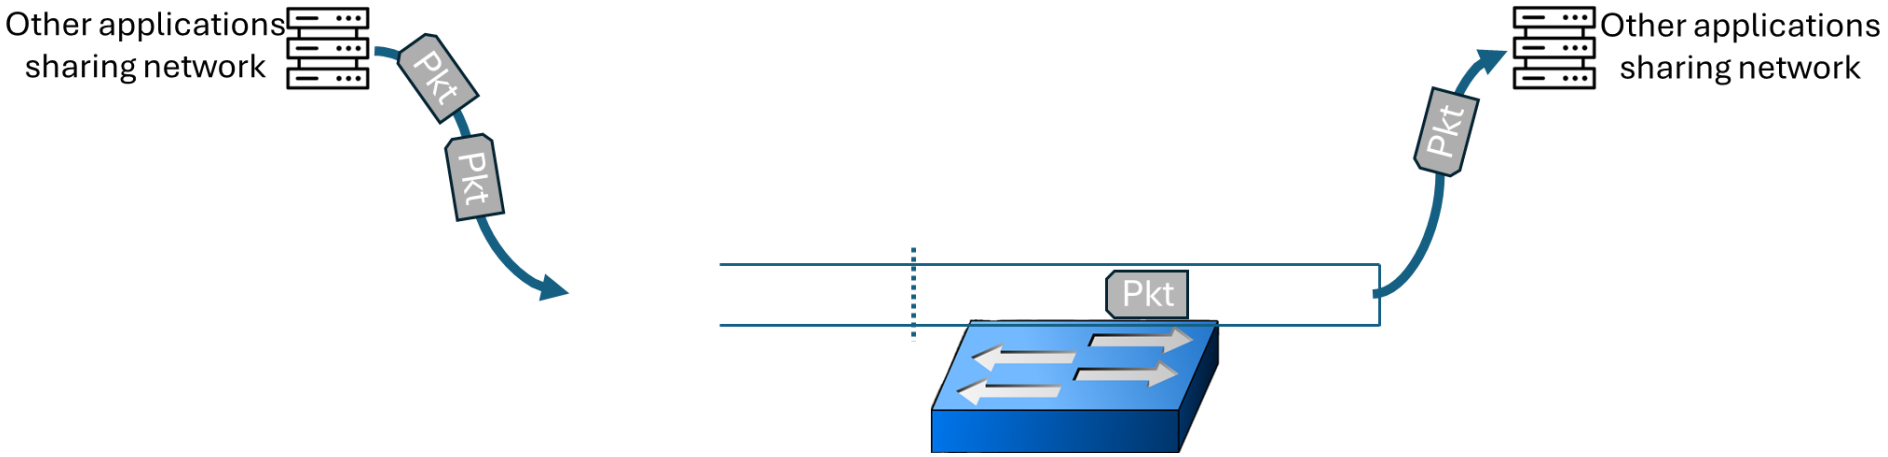

Example: Mitigating compression with just-in-time compression (HotNets 2024)

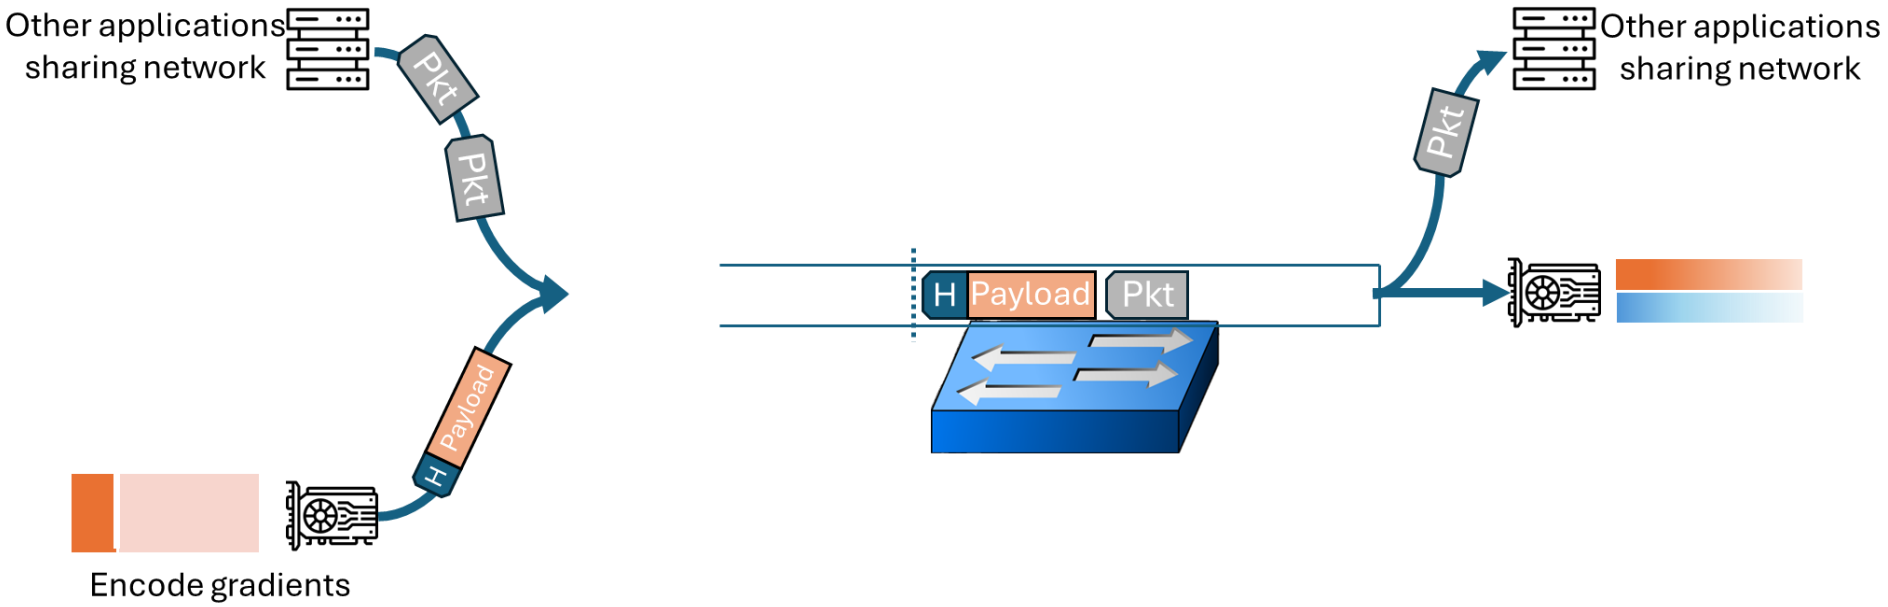

Example: Mitigating compression with just-in-time compression (HotNets 2024)

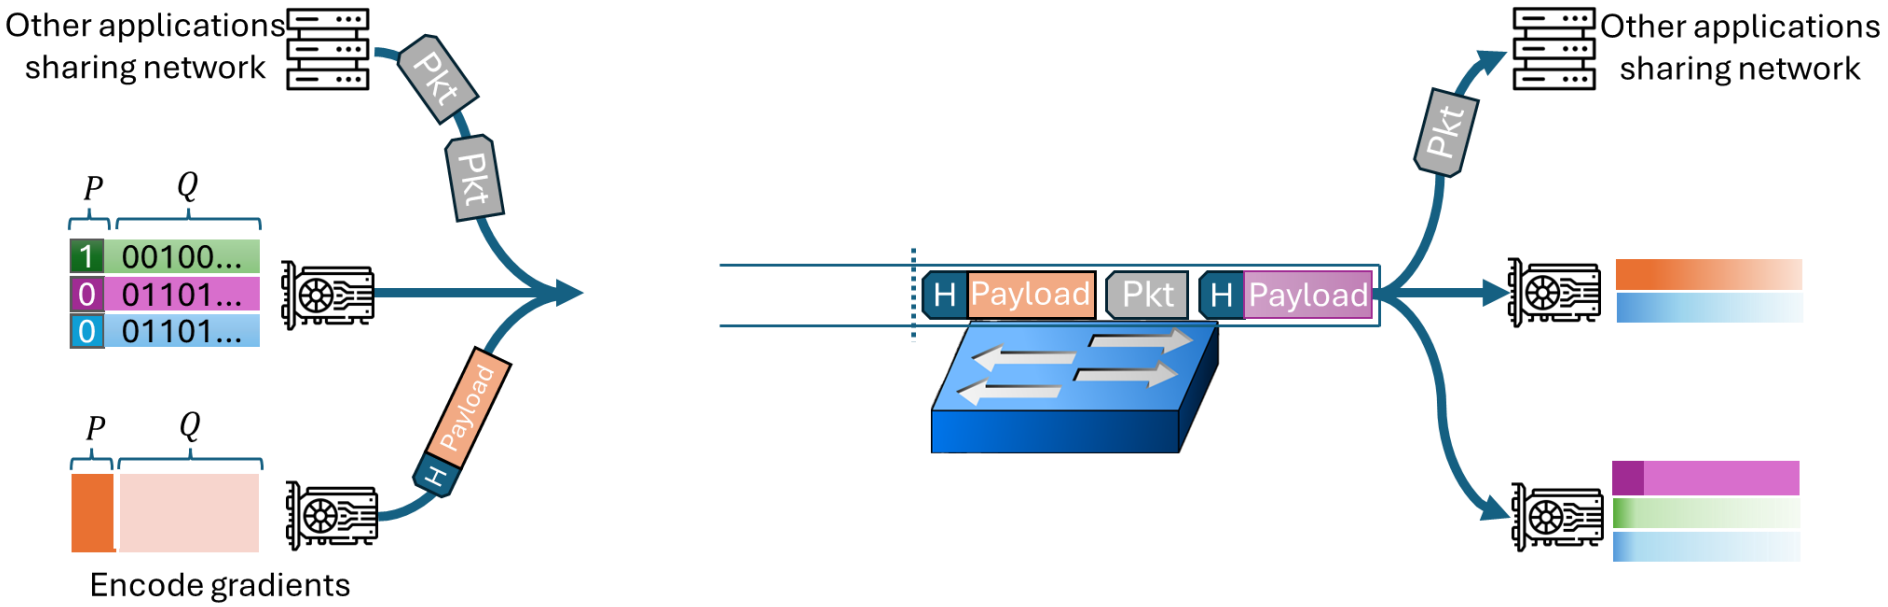

Example: Mitigating compression with just-in-time compression (HotNets 2024)

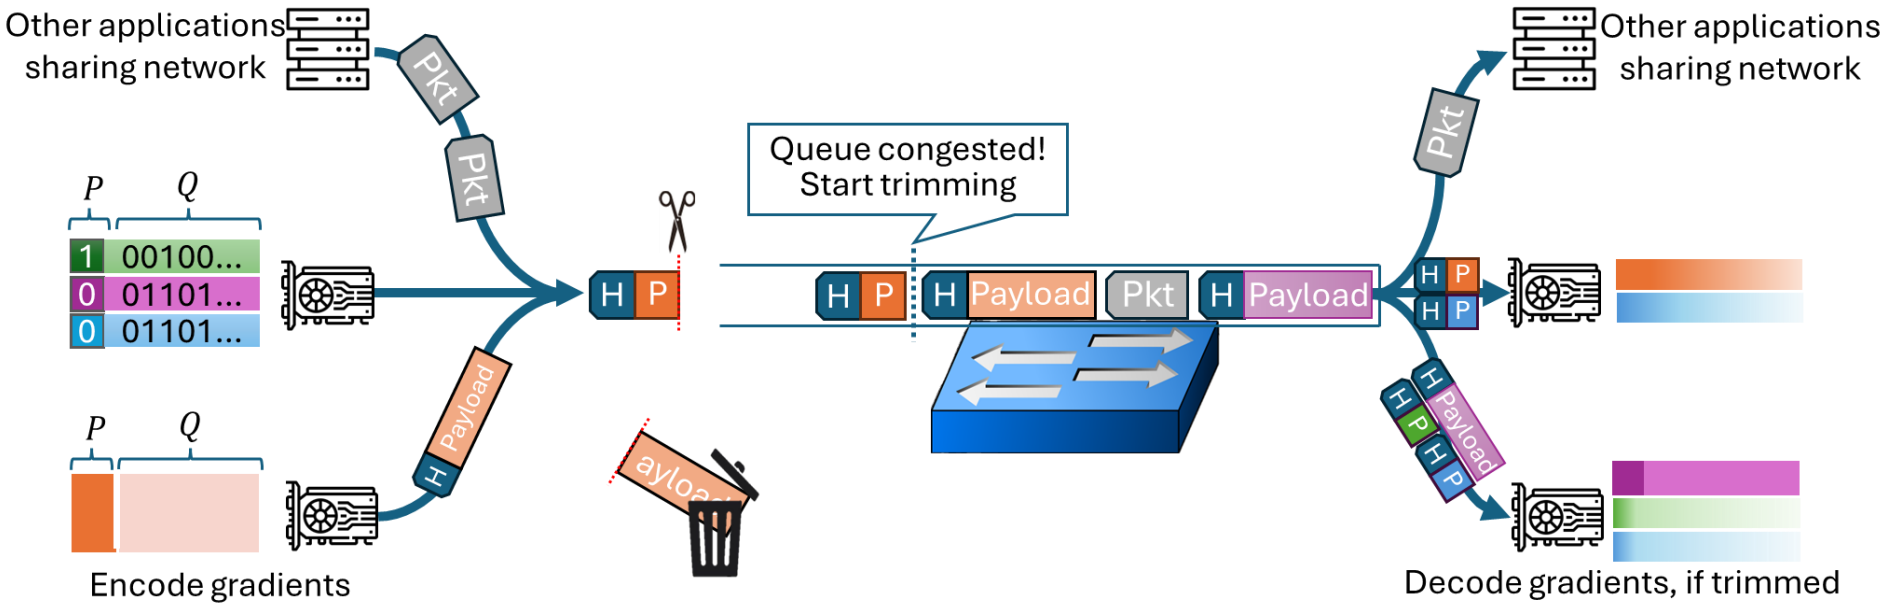

Example: Mitigating compression with just-in-time compression (HotNets 2024)

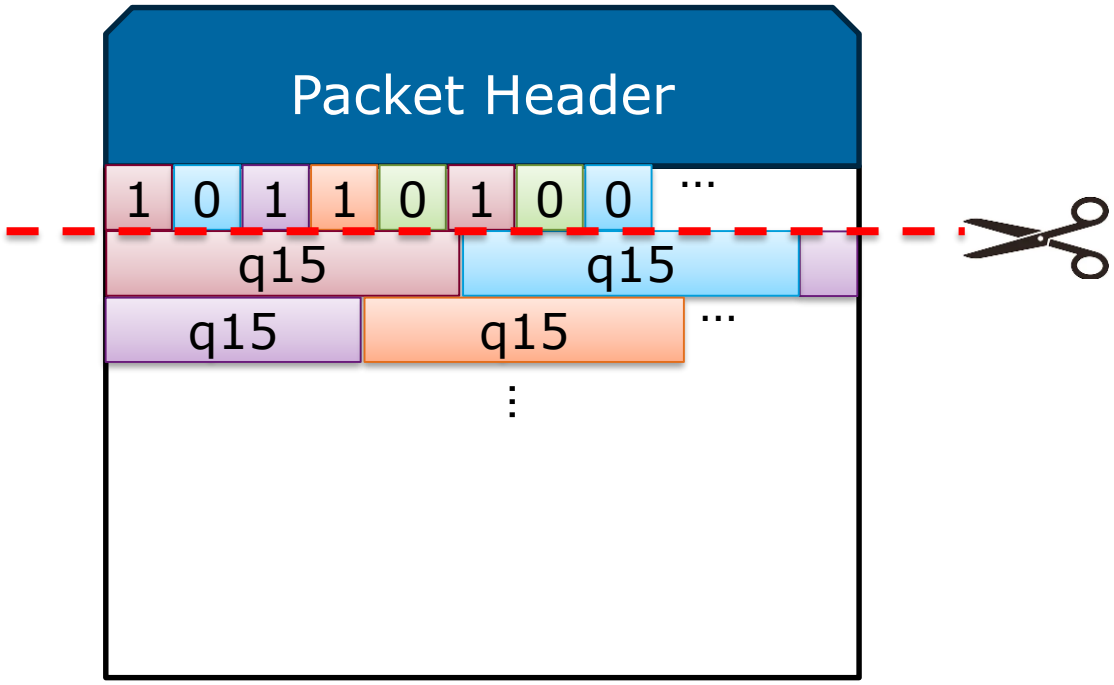

- The one-bit that makes it serves as state-of-the-art compression algorithm! (DRIVE, NeurIPS 2021).

Reproducibility

- Lossy and stochastic synchronization does **not** mean non-reproducible results!
- By logging the sources of loss, we can replay the execution of the process. For example, we log:
  - Which stragglers were dropped.
  - Which packets were trimmed.
  - …

Vision

- We develop an understanding of how system optimizations affect gradient variance and thus the convergence rate.
- We can trade a small (e.g., 5%) increase in either the #rounds or batch size into a “budget” that the system optimizations can leverage.
- We measure the benefit in **time** taken to train a model.

Example (THC, NSDI 2024)

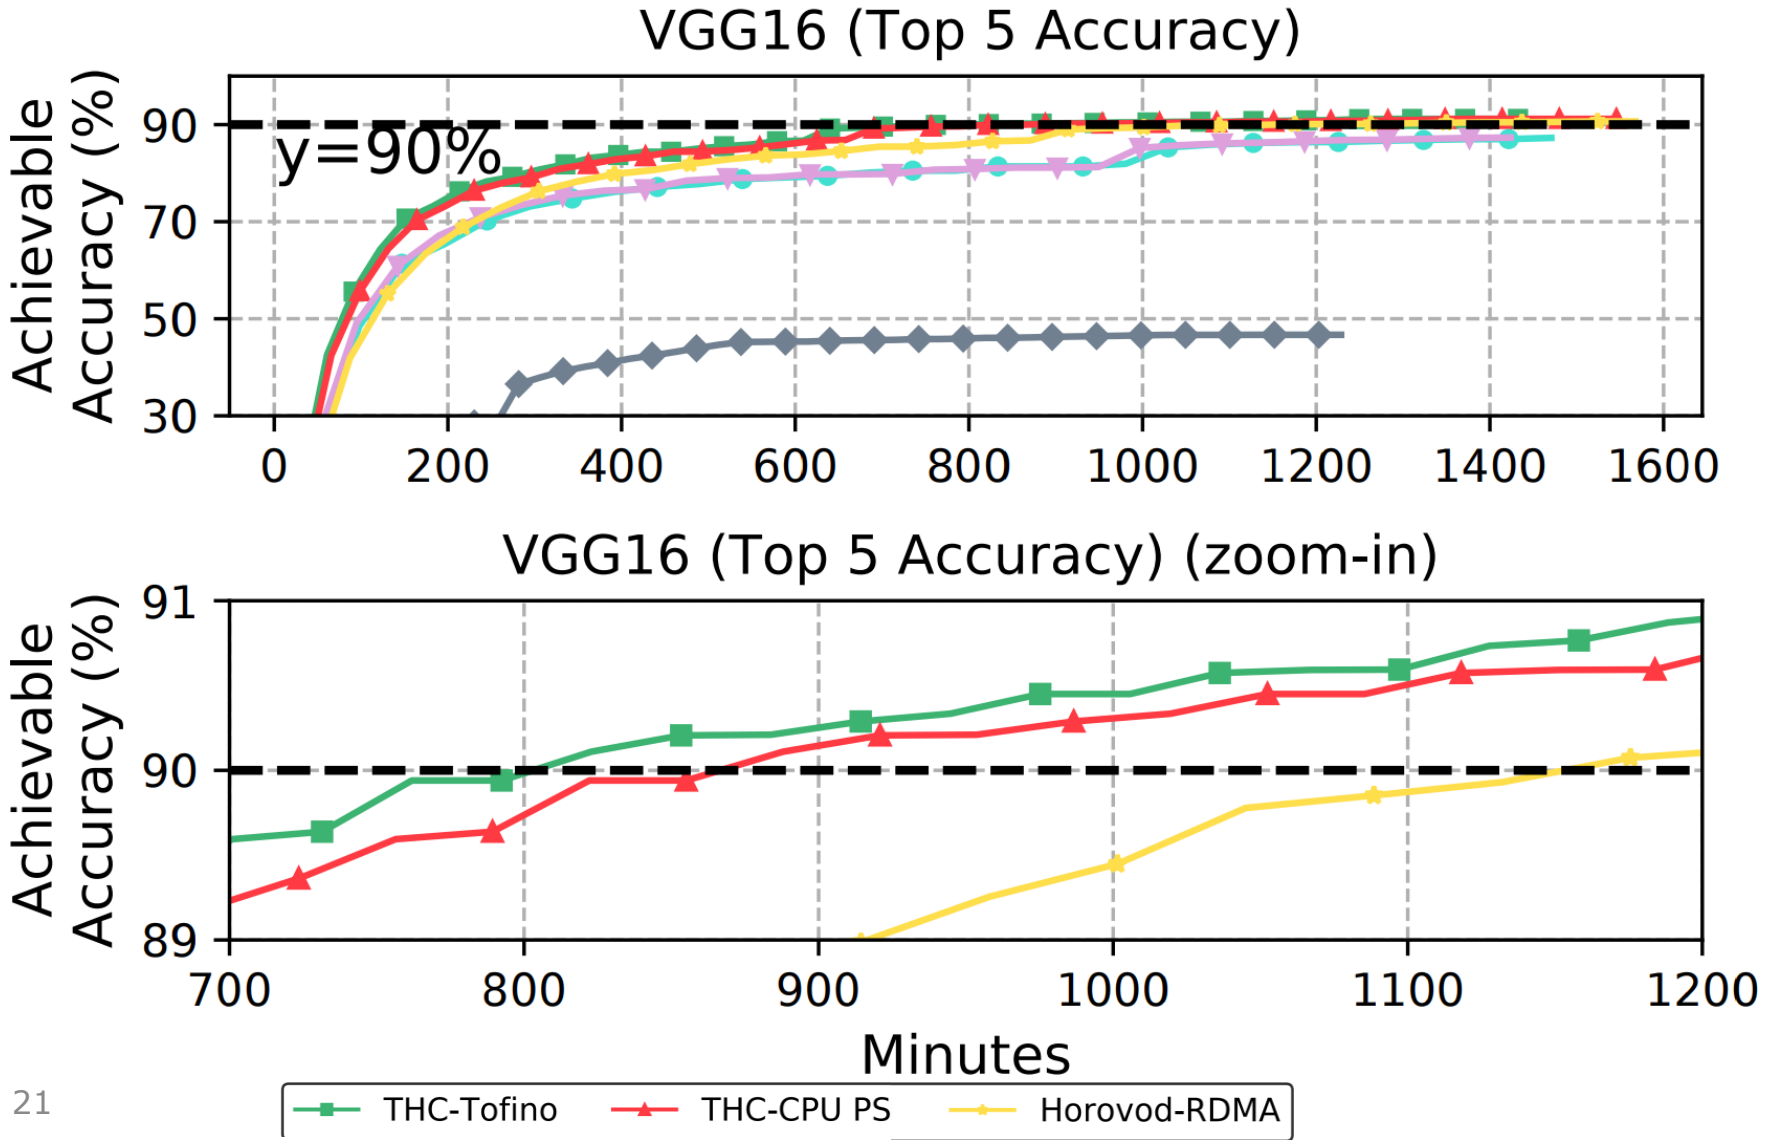

Summary

LOSE THE

LOSSLESS

ABSTRACTION

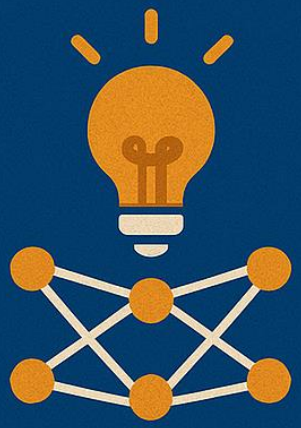

Collaborators and funders